

BIRDS-OF-A-FEATHER

AI at UFIT Research Computing



Fundamentals of Accelerated Neural Network Training with Multi-GPUs on HiPerGator-AI

Yunchao Yang, PhD

AI support team

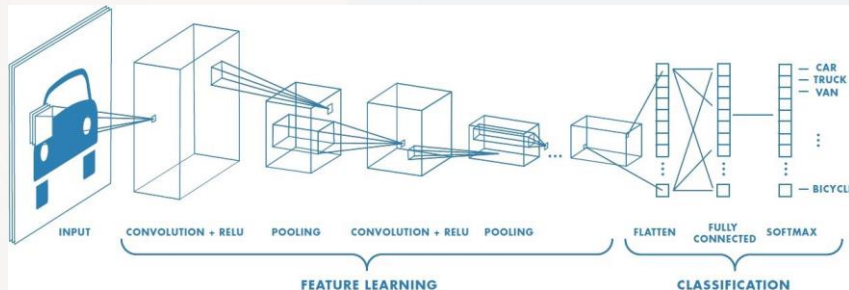
UF Research Computing

yunchyang@ufl.edu

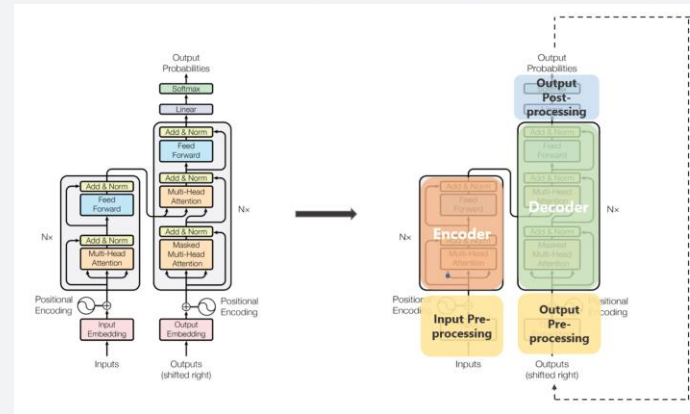
The slides and demo code will be hosted on
<https://github.com/YunchaoYang/BoF-MultiGPUTutorial>

Deep Neural Networks is transforming everywhere


Convolutional Neural Networks In Computer Vision



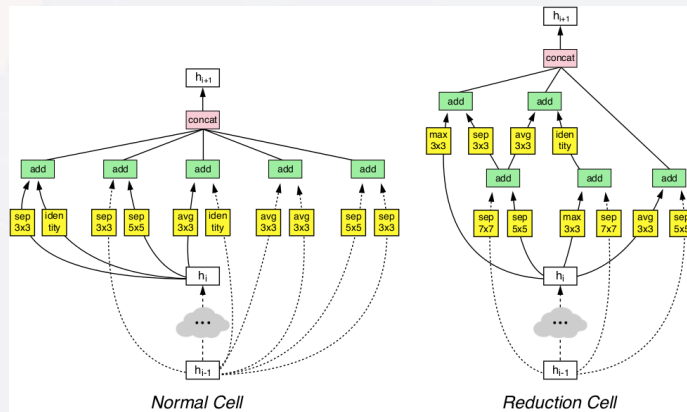
Transformers in NLP




Physics-informed Machine Learning



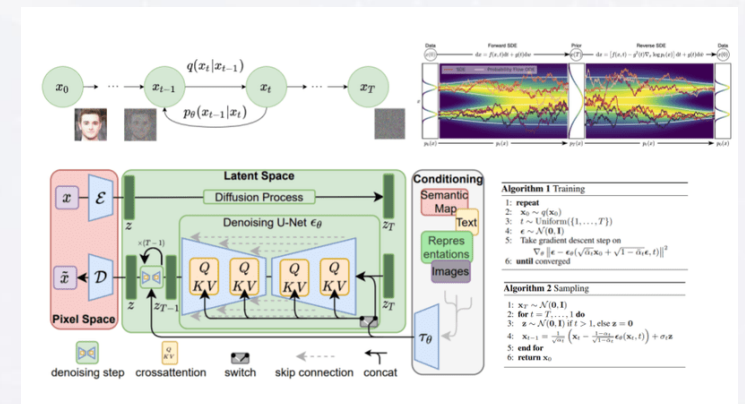
Neural Architecture Search



Reinforcement Learning



Diffusion models



Challenges for accelerating neural networks training


Neural networks are getting bigger and bigger to trillions of parameters.

GPT-3 has **175,000,000,000** parameters, at least **800 GB** memory to load the model

- Memory requirement
 - Huge parameter space
 - large batch size
- Computationally intensive
 - High FLOPs on training
 - Massive amount of training dataset (millions of images)
 - Large number of parameters – (GPT-3, 175 billion trainable weights)
 - Inference speed

Large Neural Networks

- Deep learning models are getting bigger and bigger:
 - **GPT-3: 175 billion parameters**
 - **Megatron-Turing: 530 billion parameters**



| Model | # parameters | Memory requirement for parameters (FP32) | type |
|-----------------------|--------------|--|-----------------------------|
| BERT-L | 345 m | ~ 2 GB | NLP, transformer |
| GPT-3 | 175 b | ~ 800 GB | NLP, transformer |
| Megatron-Turing | 530 b | ~ 2 TB | Natural Language Generation |
| SEER(SElf-SupERvised) | 10 b | ~ 40 GB | computer vision |

↓ can only scale using [model parallel](#) training

single GPU memory on HPG-AI: an A100 has 80 GB vRAM

CUDA Out Of Memory (OOM) error

```
RuntimeError: CUDA out of memory. Tried to allocate 128.00 MiB
```

- OOM (Out Of Memory) errors can occur when building and training a neural network model on the GPU. The size of the model is limited by the available memory on the GPU.

| | vRAM | CUDA® Cores |
|----------------------|-------|-------------|
| A100 Tensor Core GPU | 80 GB | 7936 |

| Model | # parameters | Memory requirement for parameters (FP32) | type |
|-------|--------------|--|------------------|
| GPT-3 | 175 b | ~ 800 GB | NLP, transformer |

How to train faster

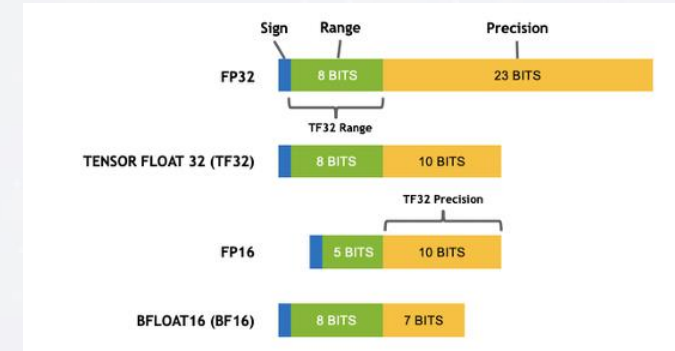


- Hardware:

- fast connectivity between GPUs
 - intra-node: NVLink
 - inter-node: Infiniband

- Software:

- Distributed and Parallel Training
 - Data parallel
 - Model parallel
- Mixed precision
 - fp16 (amp)

- GPU-to-GPU: NVLink: A100-to-A100 peer bandwidth is 300 GB/s bi-directional
- CPU to GPU: connected by PCIe Gen4 buses each provide 31.5 Gb/s for a total of 252 Gb/s



Hardware: GPUs available in HiPerGator

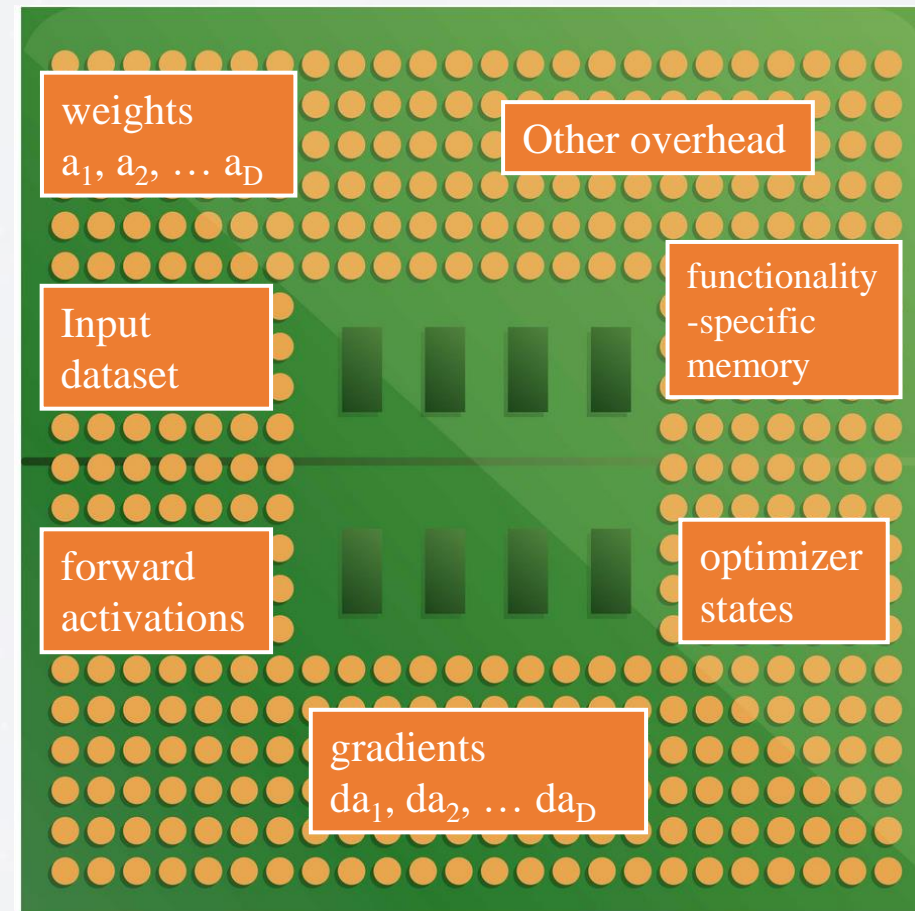
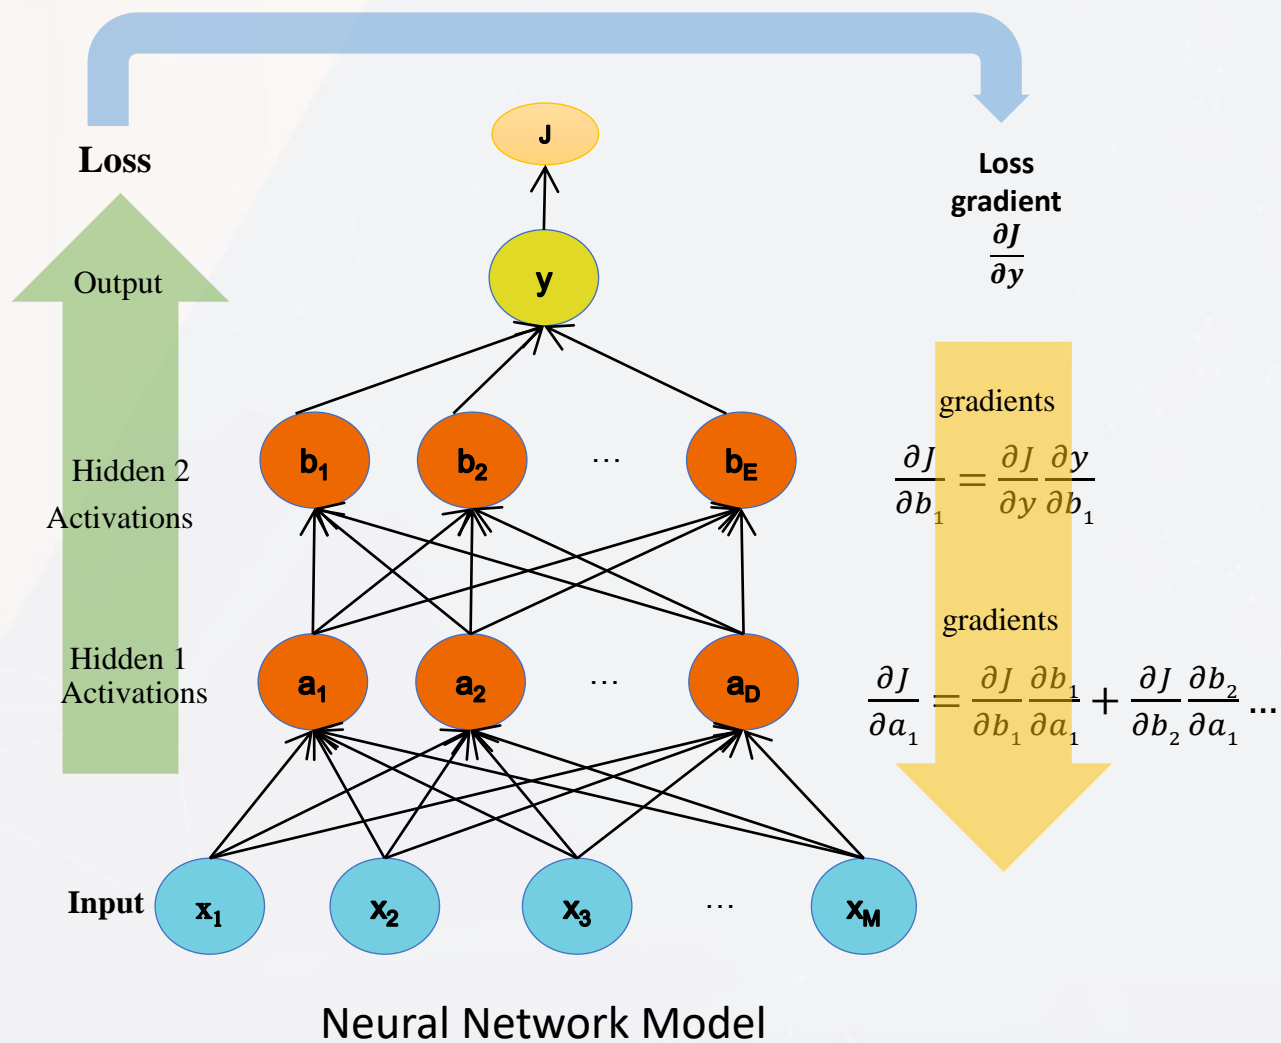
HiPerGator AI

- NVIDIA A100
 - 1,120 GPUs
 - 140 nodes
 - 8 GPUs/node
 - 80GB GPU memory per GPU

| GPU | Host Quantity | Host Architecture | Host Memory | Host Interconnect | CPUs per Host | CPUS per Socket | GPUs per Host | CPUs per GPU | Memory per GPU | SLURM Feature | GRES GPU type |
|-----------------|---------------|--------------------|-------------|-------------------|---------------|-----------------|---------------|--------------|----------------|---------------|---------------|
| GeForce 1080Ti | 1 | Intel Haswell | 128 GB | FDR IB | 28 | 14 | 2 | 14 | 11GB | n/a | geforce |
| GeForce 2080Ti | 32 | Intel Skylake | 187 GB | EDR IB | 32 | 16 | 8 | 4 | 11GB | 2080ti | geforce |
| GeForce 2080Ti | 38 | Intel Cascade Lake | 187 GB | EDR IB | 32 | 16 | 8 | 4 | 11GB | 2080ti | geforce |
| Quadro RTX 6000 | 6 | Intel Cascade Lake | 187 GB | EDR IB | 32 | 16 | 8 | 4 | 23GB | rtx6000 | quadro |
| NVIDIA A100 | 140 | AMD EPYC ROME | 2 TB | HDR IB | 128 | 16 | 8 | 16 | 80GB | a100 | a100 |

Recognize host and GPU memory are different

Anatomy of GPU memory allocation




GPU memory

Memory requirement estimation

Estimation of total VRAM memory consumption

- Consider only model parameters , gradients and optimizer states
- Ignore input, forward activations, and memory overhead
- Use FP32 data representation (4 bytes per floating number)
- Adam optimizer (storing 16 bytes per parameter)



FP32 data representation

Consider a model with one billion parameters

$$10^9 * (4B + 4B + 16B) = 10^9 * 24 \text{ Byte} = 24 \text{ GB}$$

4 bytes
per parameter


4 bytes
per gradient

16 bytes for Adam
optimizer states


$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t} \\ \hat{v}_t &= \frac{v_t}{1 - \beta_2^t} \\ \theta_{t+1} &= \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \end{aligned}$$

Fundamentals of parallel computing

Serial Computing



Parallel Computing



- A problem is broken into discrete parts that can be solved concurrently
- Each part is further broken down to a series of instructions
- Instructions from each part execute simultaneously on different processors
- An overall control/coordination mechanism is employed

Basic concepts and metrics of parallel computing

- **Computational Speed of the process**

- **Scalability**

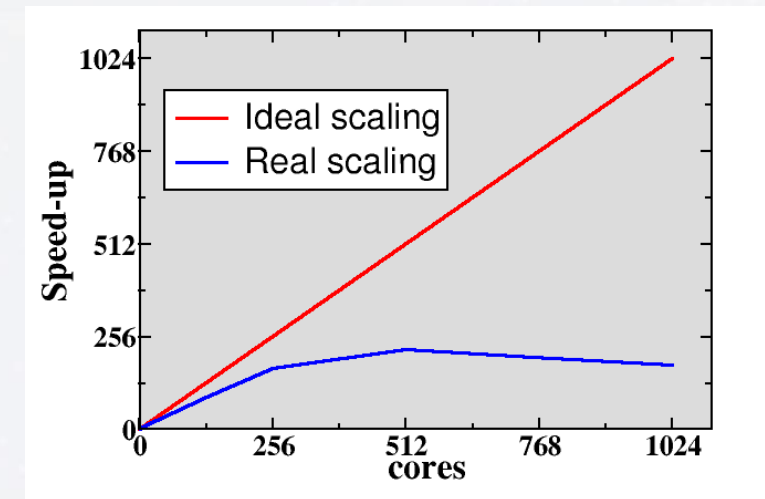
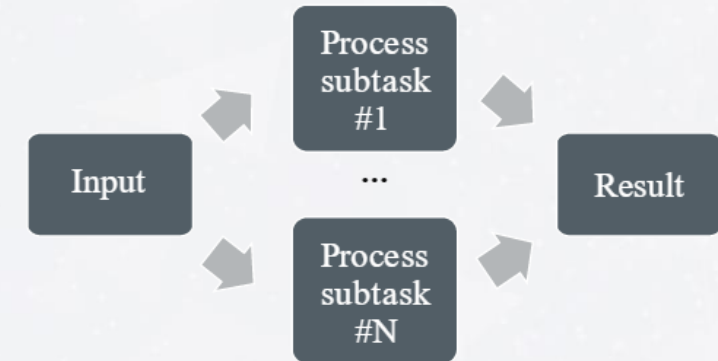
- Efficiency of a system for a growing workload

- **Speedup**

- the ratio of solution time for the sequential versus its parallel counterpart

- **Throughput**

- number of computing tasks per time unit
 - the number of images per unit time that can be processed



Types of training parallelism

- **Types of training parallelism**

- **Data Parallel**


- distributed training data for parallel optimization
 - perform allReduce on gradient

- **Model Parallelism**


- deep learning model is partitioned across multiple devices (GPU)

- Pipeline Parallel
 - Tensor Parallel
 - Optimizer-Level Parallel
 - Zero Redundancy Optimizer (ZeRO)
 - Other techniques
 - Activation checkpointing

Data Parallel

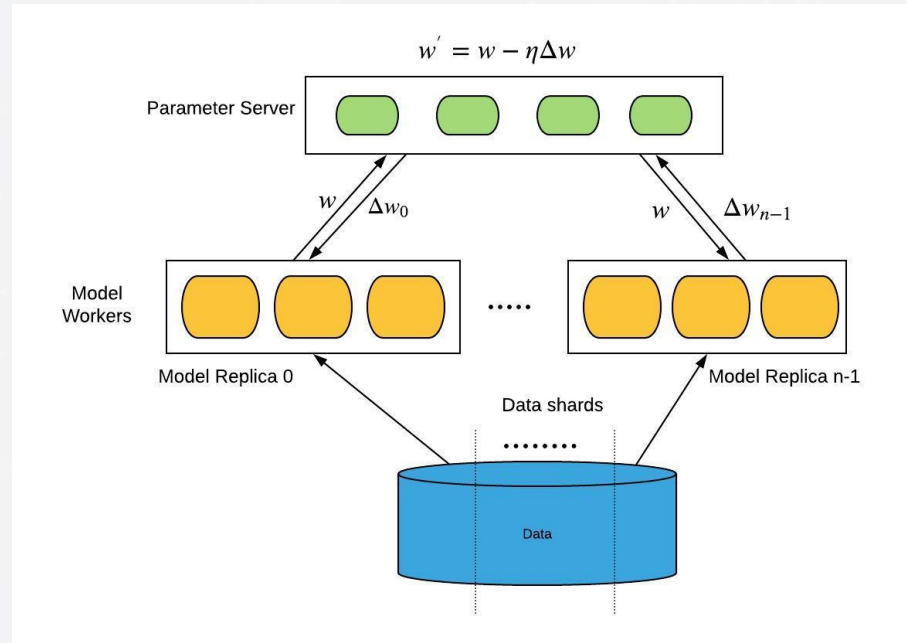


Model Parallel



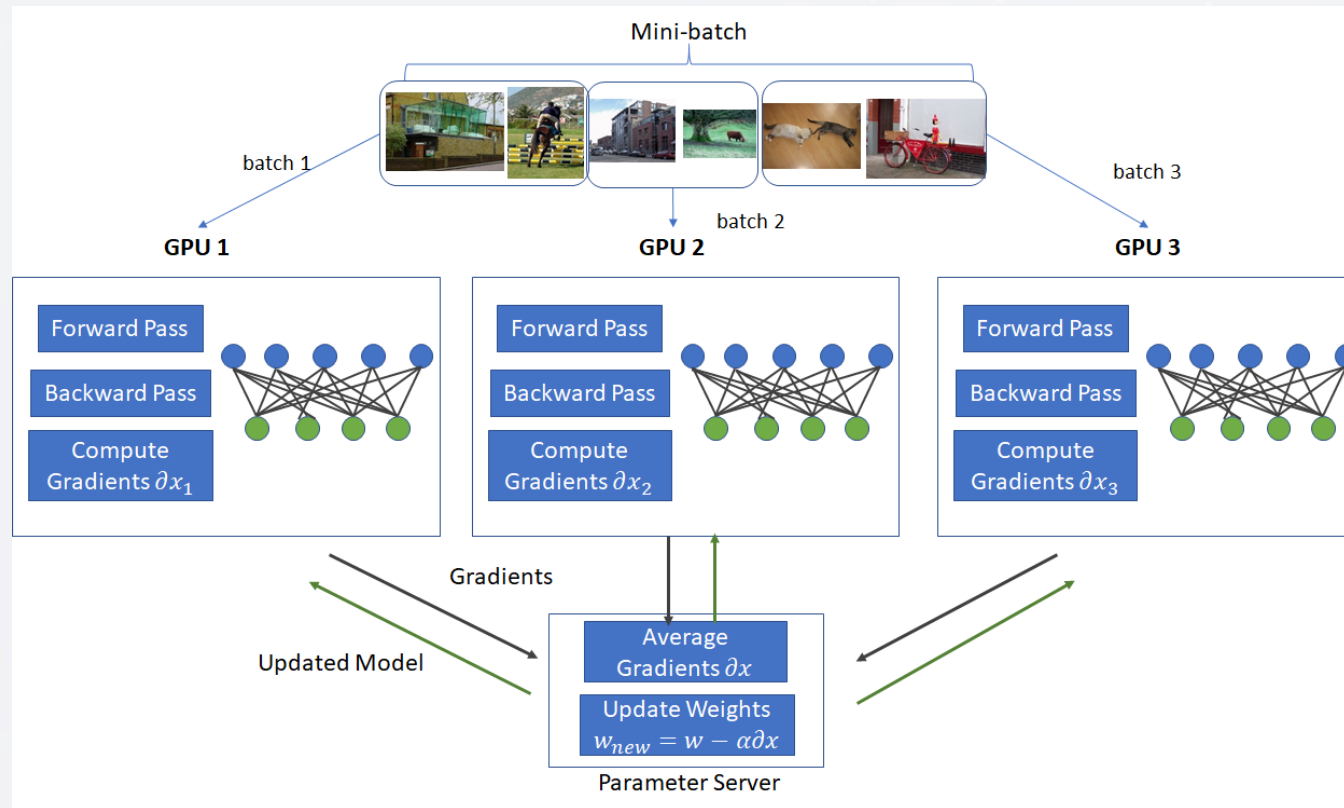
Data parallelism

- Data parallelism can be defined as the splitting of the data into N partitions where each of the partitions can be used for training into different machines or devices like CPU cores, GPUs, or even machines.
- Accumulate gradients, there are two strategies to update:
 - Synchronous Distributed
 - Asynchronous Distributed



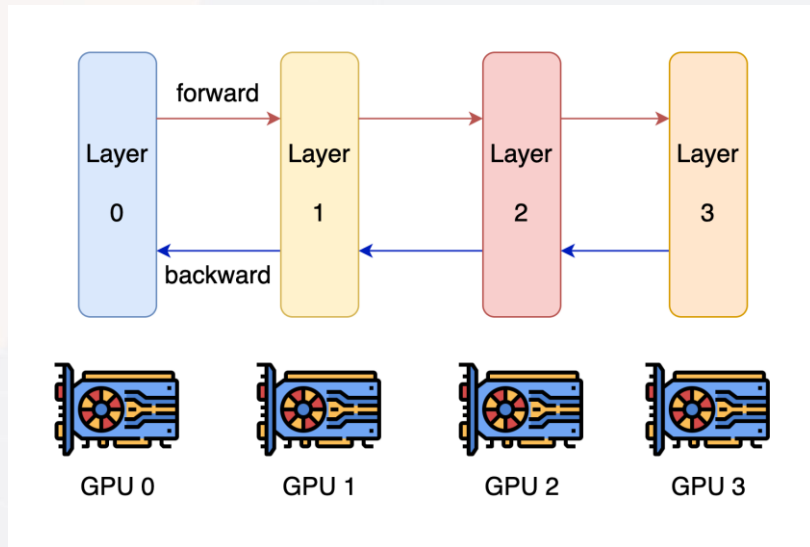
Synchronous Distributed and gradient accumulation

Divide mini-batch in micro-batches, assign on each GPU , accumulate gradient at the end of each iteration

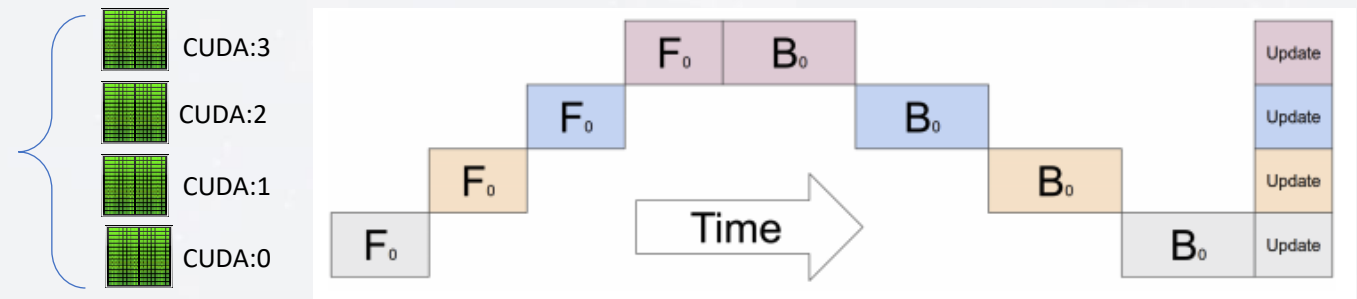


Pipeline parallelism


- Partition DNN according to depth, assign layers to specific processor



Pipeline parallel illustration



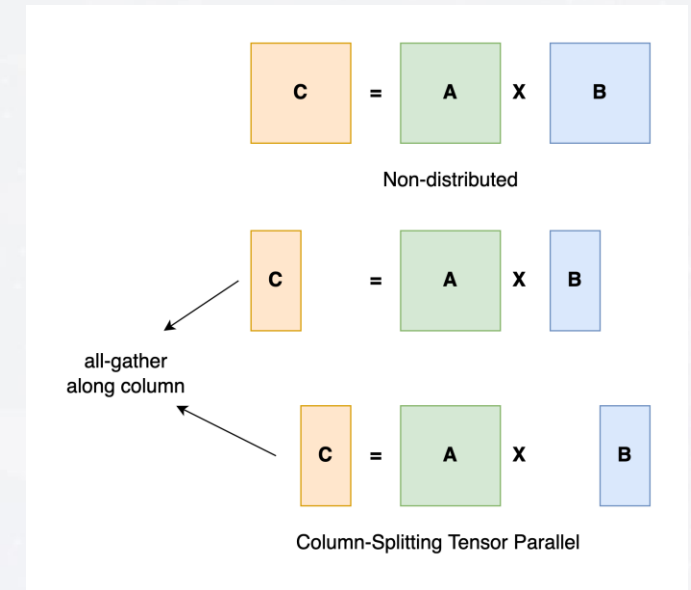
Naïve model parallelism execution: the training process suffer from GPU under utilization since only one GPU is activate.



Pipelined model parallelism Execution: splits the input minibatch into multiple microbatches and pipelines the execution of these microbatches across multiple GPUs.

Tensor parallelism

- Each tensor in one individual layer is split into multiple devices.
- each shard of the tensor resides on its designated gpu.
- Tensor parallelism is required in cases in which a single layer consumes most of the GPU memory
- Tensor parallelism is useful for extremely large models in which a pure pipelining is simply not enough. For example, in Megatron-LM library.



Tensor parallel illustration

PyTorch - Data Parallel Utilities

- PyTorch mostly provides two functions namely `nn.DataParallel` and `nn.DistributedDataParallel` to use multiple GPUs in a single node and multiple nodes during the training.
- DataParallel (DP)
 - Simple to quick start
- DistributedDataParallel (DDP)
 - robust
 - recommended

PyTorch - Data Parallel (DP)

- Implements data parallelism at the module level.

```
torch.nn.DataParallel(module, device_ids=None, output_device=None, dim=0)
```


```
>>> net = torch.nn.DataParallel(model, device_ids=[0, 1, 2])  
>>> output = net(input_var)  # input_var can be on any device, including CPU
```

- Simple to use, just wrap your model by `nn.DataParallel`, Pytorch will do everything else for you:
 - 1) replicates model to each participating GPU
 - 2) split minibatch to GPUs
 - 3) sync the gradients.
- Single process, multi-thread, subject to GIL mutex
- Only applicable to GPUs on the same node (machine).
- Does not work with model parallel
- DistributedDataParallel (DDP) is preferred instead of DataParallel (DP) even if on a single node.

Try to use PyTorch DDP instead of DP !

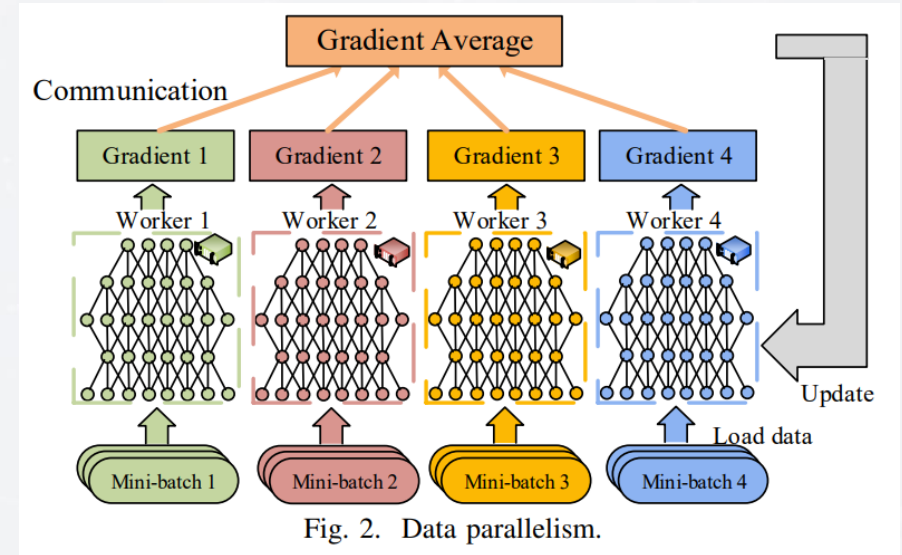
PyTorch - Distributed Data Parallel (DDP)

- Multiprocessing with DistributedDataParallel duplicates the model across multiple GPUs, each of which is controlled by one process.
- The GPUs can all be on the same node or spread across multiple nodes.
- Every process does identical task, and each process communicates with all the others.
- Only gradients are passed between the processes/GPUs so that network communication is less of a bottleneck.

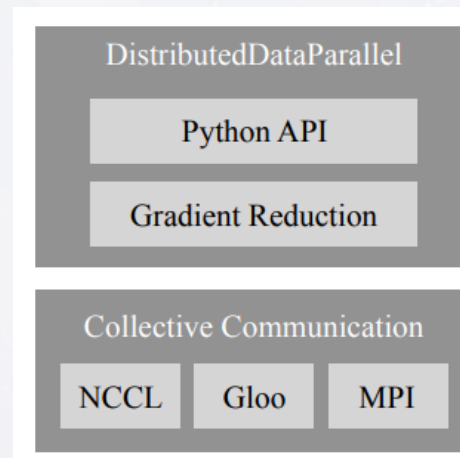


PyTorch - Distributed Data Parallel (DDP)

- [DistributedDataParallel](#) (DDP) API implements data parallelism at the module level which can run across multiple machines.
- Collective communications in the `torch.distributed` package;
- Parameters:
 - **Number of GPUs:**
 - Spawn up N processes for N GPUs.
 - Each process exclusively works on a single GPU from 0 to N-1.
 - **Batch size:**
 - The batch size should be larger than the number of GPUs;
 - **Learning step**
 - Multiple learning step by number of GPUs
 - **Collective Communication Backend**
 - nccl backend is currently the fastest and highly recommended backend when using GPUs.



Tang et al. 2020



Li et al 2019

PyTorch DDP code walkthrough

- 1. Constructing the process group
- 2. Wrap the model with DDP
- 3. Distribute dataloader (DistributedSampler wrapper)
- 4. Launch with process
 - using torchrun utility
 - torch.multiprocessing.spawn function
- 5. Clean up the process group


Pytorch DDP code walkthrough

- 1. import required libraries

```
+ import torch.multiprocessing as mp
+ from torch.utils.data.distributed import DistributedSampler
+ from torch.nn.parallel import DistributedDataParallel as DDP
+ from torch.distributed import init_process_group, destroy_process_group
+ import os
```

- 2. Constructing the process group

```
+ def ddp_setup(rank: int, world_size: int):
+     """
+     Args:
+         rank: Unique identifier of each process
+         world_size: Total number of processes
+     """
+     os.environ["MASTER_ADDR"] = "localhost"
+     os.environ["MASTER_PORT"] = "12355"
+     init_process_group(backend="nccl", rank=rank, world_size=world_size)
```



For single node only, for multi-node, use slurm to set "MASTER_ADDR"
localhost can be used for single node

Pytorch DDP code walkthrough

- 3. wrap the DDP model

```
+ self.model = DDP(model, device_ids=[gpu_id])
```

- 4. Distributing input dataloader

```
train_data = torch.utils.data.DataLoader(  
    dataset=train_dataset,  
    batch_size=32,  
    - shuffle=True,  
    + shuffle=False,  
    + sampler=DistributedSampler(train_dataset),  
)
```

- 5. Calling the `set_epoch()` method on the *DistributedSampler* at the beginning of each epoch

```
def _run_epoch(self, epoch):  
    b_sz = len(next(iter(self.train_data))[0])  
    + self.train_data.sampler.set_epoch(epoch)  
    for source, targets in self.train_data:  
        ...  
        self._run_batch(source, targets)
```


Pytorch DDP code walkthrough

- 7. launch the training process without using torchrun utility
 - 7.1 use `mp.spawn()` to spin up one process for each GPU

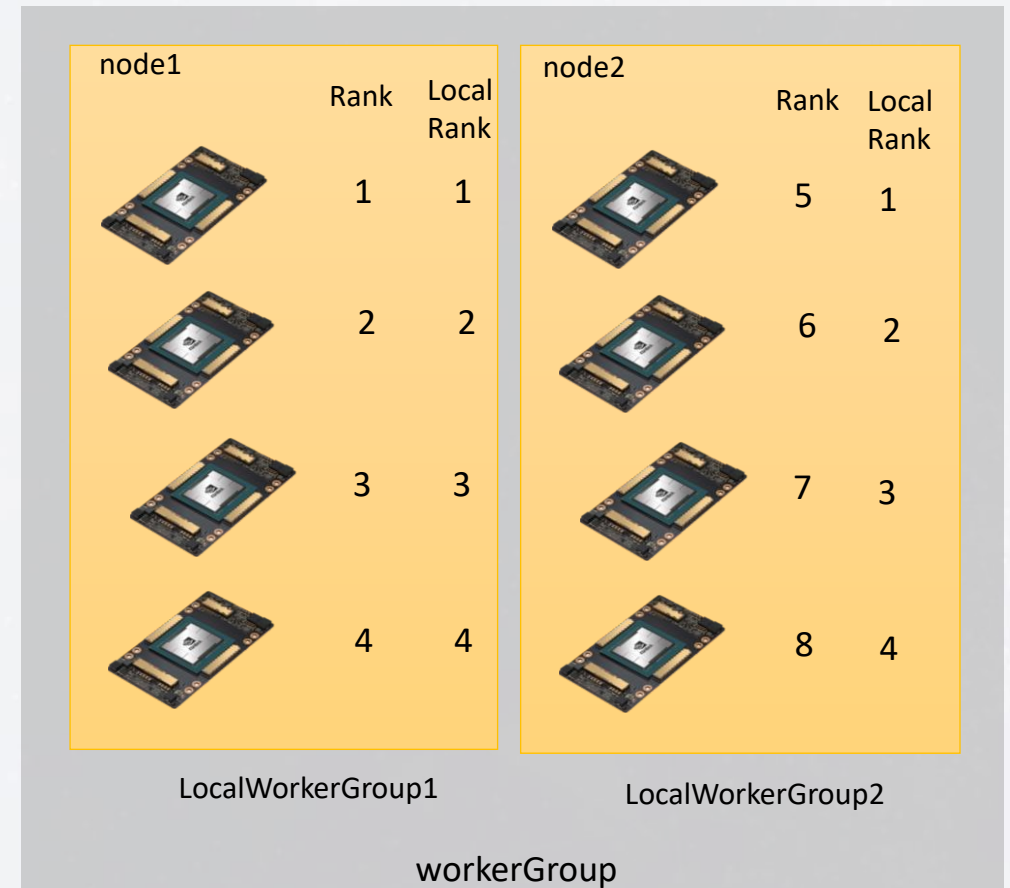
```
if __name__ == "__main__":  
    import sys  
    total_epochs = int(sys.argv[1])  
    save_every = int(sys.argv[2])  
    - device = 0      # shorthand for cuda:0  
    - main(device, total_epochs, save_every)  
    + world_size = torch.cuda.device_count()  
    + mp.spawn(main, args=(world_size, total_epochs, save_every,), nprocs=world_size)
```

- How to launch in terminal?
 - Exactly same as single GPU training:
`python train_script.py --parameters`

Torchrun utility

- torchrun is a python console script, equals to
python -m torch.distributed.run
- Definitions for torchrun
 - WORLD_SIZE
 - RANK
 - LOCAL_RANK
 - gpu_id
 - LOCAL_WORLD_SIZE
 - MASTER_ADDR
 - MASTER_PORT

RANK and WORLD_SIZE are assigned automatically by torchrun



Launch MultiGPU process - torchrun

- 7. launch the training process using torchrun
 - Single-node multi-GPU
 - Multi-node multi-GPU
 - Use torchrun to launch on each node

Changes to your training script:

```
import os
local_rank =
int(os.environ["LOCAL_RANK"])
```

```
import torch.distributed as dist
dist.init_process_group(backend="gloo|nccl")
```

launch in terminal for single Node

```
torchrun
--standalone
--nnodes=1
--nproc_per_node=$NUM_TRAINERS
YOUR_TRAINING_SCRIPT.py (--arg1 ... train script args...)
```

launch in SLURM script for single or multi node:

```
srun --export=ALL torchrun \
--nnodes 4 \
--nproc_per_node 1 \
--rdzv_id $RANDOM \
--rdzv_backend c10d \
--rdzv_endpoint $head_node_ip:29500 \
multigpu_torchrun.py 50 10
```


Demo time

- OpenOnDemand
 - 1. Launch a OpenOnDemand terminal with 1 Node and 2 GPUs
 - 2. training with single GPU
 - 3. training with 2 GPUs in 1 node
- SLURM training job
 - 4. submit a SLURM training job on multi-Node and multi-GPUs

Questions?